

Defining the Q-function

$$R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

Total reward, R_t , is the discounted sum of all rewards obtained from time t

$$Q(s_t, a_t) = \mathbb{E}[R_t | s_t, a_t]$$

The Q-function captures the **expected total future reward** an agent in state, s , can receive by executing a certain action, a

How to take actions given a Q-function?

$$Q(s_t, a_t) = \mathbb{E}[R_t | s_t, a_t]$$

↑ ↑
(state, action)



Ultimately, the agent needs a **policy** $\pi(s)$, to infer the **best action to take** at its state, s

Strategy: the policy should choose an action that maximizes future reward

$$\pi^*(s) = \operatorname{argmax}_a Q(s, a)$$



Deep Q Networks (DQN)

How can we use deep neural networks to model Q-functions?



Deep Q Networks (DQN): Training


How can we use deep neural networks to model Q-functions?



What happens if we take all the best actions?
Maximize target return → train the agent


Deep Q Networks (DQN): Training

How can we use deep neural networks to model Q-functions?




Deep Q Networks (DQN): Training

How can we use deep neural networks to model Q-functions?



Deep Q Networks (DQN): Training

How can we use deep neural networks to model Q-functions?




$$\mathcal{L} = \mathbb{E} \left[\left\| \left(r + \gamma \max_{a'} Q(s', a') \right) - Q(s, a) \right\|^2 \right]$$


Q-Loss

Deep Q Network Summary

Use NN to learn Q-function and then use to infer the optimal policy, $\pi(s)$



DQN Atari Results



Downsides of Q-learning

Complexity:

- Can model scenarios where the action space is discrete and small
- Cannot handle continuous action spaces

Flexibility:

- Policy is deterministically computed from the Q function by maximizing the reward → cannot learn stochastic policies

To address these, consider a new class of RL training algorithms:
Policy gradient methods

Deep Reinforcement Learning Algorithms

Value Learning

Find $Q(s, a)$

$a = \operatorname{argmax}_a Q(s, a)$


Policy Learning

Find $\pi(s)$

Sample $a \sim \pi(s)$

Deep Q Networks (DQN)


DQN: Approximate Q-function and use to infer the optimal policy, $\pi(s)$



Policy Gradient (PG): Key Idea

DQN: Approximate Q-function and use to infer the optimal policy, $\pi(s)$


Policy Gradient: Directly optimize the policy $\pi(s)$



What are some advantages of this formulation?

Discrete vs Continuous Action Spaces

Discrete action space: which direction should I move?




Discrete vs Continuous Action Spaces

Discrete action space: which direction should I move?




Continuous action space: how fast should I move?

0.7 m/s



Policy Gradient (PG): Key Idea

Policy Gradient: Enables modeling of continuous action space



Deep Reinforcement Learning: Summary

Foundations

- Agents acting in environment
- State-action pairs → maximize future rewards
- Discounting



Q-Learning

- Q function: expected total reward given s, a
- Policy determined by selecting action that maximizes Q function



Policy Gradients

- Learn and optimize the policy directly
- Applicable to continuous action spaces

